

# Tutorial: Training a Classifier

Friday, July 17, 2020

4:32 PM

## [Data](#)

### [For this tutorial](#)

### [Steps for Training an image classifier](#)

1. [Loading and normalizing CIFAR10](#)
2. [Define a Convolutional Neural Network](#)
3. [Define a Loss Function and Optimizer](#)
4. [Train the Network](#)
5. [Test the network on the test data](#)

## [Training on GPU](#)

### [Training on multiple GPUs](#)

---

## Data

You can use *standard python packages* that **load data into a *numpy array***

- Then you can **convert** this array into a `torch.*Tensor`
- *Python Packages:*
- For **images** => can use `Pillow` and `OpenCV`
- For **audio** => can use `scipy` and `librosa`
- For **text** => can use:
  - o either **raw Python** or **Cython based loading**
  - o or `NLTK` and `SpaCy`

For **vision**, use the **package** => `torchvision`

- It has data loaders for common datasets such as `Imagenet`, `CIFAR10`, `MNIST`, etc. and data transformers for images, visualizations
  - o `torchvision.datasets` and `torch.utils.data.DataLoader`

---

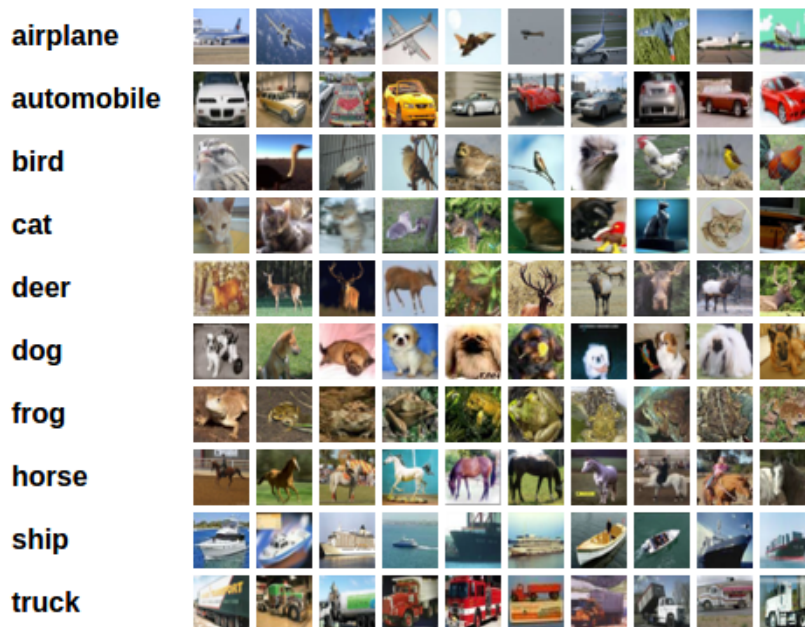
## For this tutorial

Tutorial Link:

- [https://pytorch.org/tutorials/beginner/blitz/cifar10\\_tutorial.html#sphx-glr-beginner-blitz-cifar10-tutorial-py](https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html#sphx-glr-beginner-blitz-cifar10-tutorial-py)

Uses the **CIFAR10** dataset:

- It has the **classes**: 'airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck'.
- The **images in CIFAR-10** are of **size 3x32x32**
  - o **Example:**
    - **3-channel color** images of **32x32 pixels** in size.



---

## Steps for Training an image classifier

We will do the following steps in order:

1. **Load** and **normalizing** the **CIFAR10 training** and **test datasets** using **torchvision**
2. **Define** a **Convolutional Neural Network (CNN)**
3. **Define** a **loss** function

4. **Train** the **network** on the *training data*
5. **Test** the **network** on the *test data*

---

## 1. Loading and normalizing CIFAR10

The **output** of *torchvision datasets* are **PILImage** images of range [0, 1].

- Need to **transform them to Tensors** of normalized range [-1, 1].

### STEPS:

- a. **Create** the **transform function**.
  - This **creates tensors** and **normalizes** the data.
- b. **Get** the **training set data** from `torchvision.datasets`.
  - Set argument `train=True`
  - Set transform function to transform argument
- c. **Pass** the **training set** to the **DataLoader** using `torch.utils.data.DataLoader`.
  - This adds the `batch size` and `num_workers`, etc.
- d. **Repeat** *step b and c* for **test set**.
- e. **Create** the **list of classes** for the dataset as a *python tuple*.

---

## 2. Define a Convolutional Neural Network

- Example:
  - Takes **3 channels images**

```
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
```

```

        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()

```

### 3. Define a Loss Function and Optimizer

#### USE:

- **Loss function** => *Classification Cross-Entropy loss*
- **Optimizer** => *SGD with momentum*
  - Updates the weights
  - Need to **set a learning rate** (example: lr=0.001)
  - Need to **set momentum** (example: momentum=0.9)
    - ???

### 4. Train the Network

Model **loops** over the data iterator, and **feeds the inputs to the network and optimize**.

- Each loop is called an **epoch**
- **GOAL** is to **lower the loss**.
  - By **mini batches**. (example: loss for every 2000 samples)
- Then **save** the **trained model**.
  - More details on saving PyTorch models:
    - <https://pytorch.org/docs/stable/notes/serialization.html>
  - Example:

```
PATH = './cifar_net.pth'
torch.save(net.state_dict(), PATH)
```

---

## 5. Test the network on the test data

Need to **check if the network has learnt** anything at all.

- A. **By predicting the class label** that the neural network **outputs**,
  - B. and **by checking it against the ground-truth** (*test dataset*)
- If the prediction is correct, we ***add the sample to the list of correct predictions.***

### STEPS:

1. Check and **view a few images** and **corresponding labels** from the *test dataset*.
2. **Load** back **saved model**.
  - *Note:*
    - **Saving and re-loading the model** is only necessary if it's going to be open elsewhere:

```
net = Net()
net.load_state_dict(torch.load(PATH))
```
3. **Run testing images** through model.
4. **Get the index of the highest energy:**
  - The **outputs are *energies*** for the 10 classes.
  - The higher the energy for a class, the more the network thinks that the image is of the particular class.
5. **Look at how** the network performs on the **whole dataset** by checking the ***accuracy***.
6. **Check** the performance (***accuracy***) of every class.

---

## Training on GPU

Just like how you transfer a Tensor onto the GPU, you *transfer the neural net onto the GPU*.

- First **define our device** as the *first visible cuda device* if we have CUDA available:

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
# Assuming that we are on a CUDA machine, this should print a CUDA device:
print(device)
```

➤ The rest of this section assumes that `device` is a CUDA device.

- Then these **methods** will recursively go *over all modules* and **convert** their **parameters** and **buffers** to CUDA tensors:

```
net.to(device)
```

- You **will have to send the inputs and targets at every step** to the GPU too:

```
inputs, labels = data[0].to(device), data[1].to(device)
```

- To increase the speedup on the model:
  - Try increasing the width of the network.
    - By changing **argument 2** of the **first** `nn.Conv2d`,
    - and **argument 1** of the **second** `nn.Conv2d`,
    - they need to be the same number!

---

## Training on multiple GPUs

For **more even more speedup** using all of your GPUs, please check out:

- doc: `data_parallel_tutorial`