

# PyTorch Fundamentals

Thursday, July 9, 2020

11:44 PM

[PyTorch Basics](#)

[Autograd: Automatic Differentiation](#)

[Neural Networks](#)

[Defining a neural network & Processing inputs and calling backward](#)

[Loss Function](#)

[Backprop](#)

[Update the weights](#)

\*Review from: [https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)

---

## PyTorch Basics

### What is PyTorch?

- It's a **Python-based** scientific **computing package** for audiences:
  - o A replacement for NumPy to **use the power of GPUs**
  - o a deep learning research platform that provides maximum flexibility and speed
- **Tensors** are similar to NumPy's ndarrays.
  - o They can also be used on a GPU to accelerate computing.
  - o Tensor operation=> <https://pytorch.org/docs/torch>
  - o **Work like a matrix / multi-dimensional arrays.**
  - o A tensor is **easily** converted to Numpy array
    - changing one will change the other because they share memory locations (if the Torch Tensor is on CPU)
    - **Except a CharTensor**
- o Any operation that mutates a tensor **in-place** is post-fixed with an `__`.
  - o Example: `x.copy_(y)`

---

# Autograd: Automatic Differentiation

The `autograd` package provides *automatic* differentiation for all **operations** on Tensors.

- Generally, `torch.autograd` is an engine for **computing vector-Jacobian product**.
- The **vector-Jacobian product** makes it very convenient to feed external gradients into a model that has non-scalar output.
- Document about `autograd.Function` is at <https://pytorch.org/docs/stable/autograd.html#function>

➤ `torch.Tensor` is the central class of the package.

- **Track all operations on it:**
  - Set its attribute `.requires_grad` as **True**
- **Have all the gradients computed automatically:**
  - When you finish your computation you can call `.backward()`
    - Called **backprop**
    - Creates the **gradient** for a tensor by:
      - ◆ Combining all equations that were computed to get to tensor and takes the derivative with respect to the initial value as the variable.
- **Gradient for this tensor will be accumulated into this attribute:**
  - `.grad`
    - *This is also the derivative equation created by `.backward()`*
    - Can be used as a function on other tensors.
- **Stop a tensor from tracking history / future tracking:**
  - You can call `.detach()`
  - Or wrap the code block in `with torch.no_grad():`
    - *2nd option is helpful if model may have trainable parameters with no gradients.*

- **Function** another very important class for autograd implementation.
    - **Tensor** and **Function** are interconnected and *build up an acyclic graph*, that **encodes a complete history of computation**.
    - Each tensor has a `.grad_fn` attribute that references a **Function** that has created the Tensor.
      - (except for Tensors created by the user - their `grad_fn` is None).
  - If you want **to compute the derivatives**, you can call `.backward()` on a **Tensor**.
    - If Tensor is a scalar (i.e. it holds a one element data).
      - You **don't need** to specify any arguments to `.backward()`
    - If Tensor **is not** a scalar (**has more elements**):
      - You **need** to specify a **gradient** argument that is a tensor of matching shape.
      - Pass a **vector (tensor)** as argument to `.backward` to compute the vector-Jacobian product
- 

## Neural Networks

Neural networks can be constructed using the **torch.nn** package.

- **nn** depends on **autograd** to *define models and differentiate them*.
  - An **nn.Module** contains:
    - **layers**,
    - and a **method** `forward(input)` that *returns the output*.
  - Full list of **modules and loss functions** that form the building blocks of deep neural networks:
    - <https://pytorch.org/docs/stable/nn.html>
- Typical training procedure for a neural network:
- **Define** the neural network that has some *learnable parameters* (or **weights**).
  - **Iterate** over a dataset of inputs.
  - **Process** input through the network.
  - **Compute the loss** (*how far is the output from being correct*).

- **Propagate gradients** back into the network's parameters.
- **Update** the **weights** of the network.
  - Typically using a simple **update rule**:

$$\text{weight} = \text{weight} - \text{learning\_rate} * \text{gradient}$$

## Defining a neural network & Processing inputs and calling backward

- **You just have to define the forward function.**
  - The **backward** function (*where gradients are computed*) is automatically defined for you using **autograd**.
  - You can use any of the **Tensor operations** in the **forward** function.
- Get the **learnable parameters** (weights) of a model by using **net.parameters()**
- Place the **input** data to model.
- **Zero out** the **gradients** before backpropagating using **.zero\_grad()**
- **Backpropagate** using the **.backward** function
- **torch.nn** only supports **mini-batches**.
  - The entire ``**torch.nn**`` package **only supports inputs that are a mini-batch of samples**, and **not a single sample**.
  - **\*\*If you have a single sample, just use ``input.unsqueeze(0)`` to add a fake batch dimension.**
- **torch.Tensor** - A multi-dimensional array with support for autograd operations like **backward()**. Also holds the gradient w.r.t. the tensor.
- **nn.Module** - Neural network module. Convenient way of encapsulating parameters, with helpers for moving them to GPU, exporting, loading, etc.
- **nn.Parameter** - A kind of Tensor, that is automatically registered as a parameter when assigned as an attribute to a **Module**.
- **autograd.Function** - Implements forward and backward definitions of an autograd operation. Every **Tensor** operation creates at least a single **Function** node that connects to functions that created a **Tensor** and encodes its history.

---

## Loss Function

A **loss function** takes the (**output**, **target**) pair as inputs, and computes a value that **estimates** how far away the output is from the target.

- There are several different loss functions:
  - <<https://pytorch.org/docs/stable/nn.html#loss-functions>>
  - Example:
    - `nn.MSELoss` which computes the *mean-squared error* between the **input** and the **target**.

```
output = net(input)
criterion = nn.MSELoss()
#make target same shape as output
loss = criterion(output, target)
```

- You can follow **loss** in the backward direction, using its `.grad_fn` attribute.

```
input -> conv2d -> relu -> maxpool2d -> conv2d -> relu -> maxpool2d
      -> view -> linear -> relu -> linear -> relu -> linear
      -> MSELoss
      -> loss
```

---

## Backprop

To **backpropagate** the error all we have to do is to `loss.backward()`.

- You need to clear the existing gradients though, **else gradients will be accumulated to existing gradients**.
  - `model.zero_grad()`
- Then call **function to backpropagate** on the **loss function**:
  - `loss.backward()`

## Update the weights

- The **weights are updated** using an **optimizer** (*update rule*).
- The simplest **update rule** used in practice is the **Stochastic Gradient Descent (SGD)**:

$$\text{weight} = \text{weight} - \text{learning\_rate} * \text{gradient}$$

- Example: (Implementation in Python.)

```
learning_rate = 0.01
for f in net.parameters():
    f.data.sub_(f.grad.data * learning_rate)
```

- There are **various different update rules** such as **SGD**, **Nesterov-SGD**, **Adam**, **RMSProp**, etc.
  - **Enable** these methods with this package: **torch.optim**
- Incorporate into the training loop.
  - Example:

```
import torch.optim as optim
# create your optimizer
optimizer = optim.SGD(net.parameters(), lr=0.01)
# in your training loop:
optimizer.zero_grad() # zero the gradient buffers
output = net(input)
loss = criterion(output, target)
loss.backward()
optimizer.step() # Does the update
```